

# Week2 CNN Architectures

**Tutor:**

**Email:**


**Tutorial:**

**Code:** <https://github.com/Jinxu-Lin/COMP5329>

# Overfitting


# Overfitting

- Reasons:
  - Training data is too small to represent the whole dataset
  - Large amount of irrelevant information
  - Model complexity is too high
  - The model is trained for too long



# Overfitting

- Solutions:
  - Increase training data (collect more data)
  - Data cleaning jobs
  - Reduce model complexity
  - Regularization (reduce the effect of model complexity)
  - Stop training early



# Techniques

# L1/L2 Regularization

- Given loss  $\hat{L}(\theta)$ , Training objective:
  - $\min_{\theta} \hat{L}(\theta) = \frac{1}{n} \sum_{i=1}^n l(\theta, x_i, y_i)$ , s.t.  $R(\theta) \leq r$ , where  $l$  is the loss for each instance.
- L2 regularization:  $R(\theta) = ||\theta||^2$
- L1 regularization:  $R(\theta) = ||\theta||$
- Bayesian regularization can actually be transformed to L2/L1 regularization. Refer to this link to see: <https://towardsdatascience.com/a-bayesian-take-on-model-regularization-9356116b6457>

# Soft Regularization

- Use regularization as penalty, we have the new objective

- $$\min_{\theta} = \frac{1}{n} \sum_{i=1}^n l(\theta, x_i, y_i) + \lambda * R(\theta)$$


- However, the calculation of  $R(\theta)$  is still very complicated in case of L2 regularization, we adopt the weight decay:

- $$\theta = \theta - w_d * \theta$$

- Normally, weight decay  $w_d \rightarrow 0$ . We can change this value during training

# Dropout


- The basic idea is to add some random/surprising factors into the training process (only training/ deactivated during testing)
  - $o = f(W * x + b)$
- $o * \text{mask}$ , half of the mask to be 0.0, half of the mask to be 1





# Dropout

- During the test time, you need to scale down or scale up in the training
- Behaviors of dropout during training and testing phase is different:
  - Magnitude of values
  - don't filter out values during testing.



# Batch Normalization

- The previous regularization only helps to shrink/cut-off redundant information/ provide surprising factors to the training process.
- We need a more essential method to tackle the difference between training data and testing data.

# Batch Normalization

- Training

- Update mean value  $\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$
- Update variance value  $\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$
- Update running mean value  $\mu_x \leftarrow E[\mu_B]$  (implementation)
- Update running var value  $\sigma_x^2 \leftarrow E[\sigma_B^2]$  (implementation)
- Calculate Normalized value  $\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$
- Update new neural value  $x_i \leftarrow \gamma \hat{x}_i + \beta$

# Batch Normalization

- Testing

- Calculate Normalized value  $\hat{x}_i \leftarrow \frac{x_i - \mu_x}{\sqrt{\sigma_x^2 + \epsilon}}$
- Update new neural value  $x_i \leftarrow \gamma \hat{x}_i + \beta$
- with  $\gamma, \beta$  are learned values,  $\mu_B, \sigma_B^2$  are accumulated during training.

# Exam-Style Questions

# Code

# Code

- `./materials/Week7_Regularization/Week7_Regularization.ipynb`
- `./ResNet/Models/batch_normalization.py`
- `./ResNet/Models/dropout.py`
- `./ResNet/Models/resnet.py`
- `./ResNet/Optimizers/sgd.py: line26-28`